

FAANG Data Scientist Resume — 2026 Template

ATS-optimised · Current tech stack · Hired at top companies · 0–4 years experience

 Customisation Guide

 Download PDF

Your Full Name

DATA SCIENTIST · MACHINE LEARNING · GENAI · BIG DATA

 City, State |  (555) 000-0000 | yourname@email.com | linkedin.com/in/yourname | github.com/yourname | yourname.dev (portfolio)

TECHNICAL SKILLS

Languages	Python (NumPy, Pandas, Scikit-learn) · SQL (PostgreSQL, BigQuery, Snowflake) · R · Scala (basics)
ML / AI	XGBoost, LightGBM, Random Forest, Logistic Regression · PyTorch · HuggingFace Transformers · Scikit-learn Pipelines
GenAI / LLM	Gemini API, OpenAI API · LangChain, LlamaIndex · RAG Pipelines · ChromaDB / FAISS / Pinecone · Prompt Engineering
Big Data	Apache Spark, PySpark · Kafka · Apache Airflow · Delta Lake · dbt · Hadoop HDFS
Cloud & MLOps	AWS (SageMaker, S3, EC2) · GCP (BigQuery, Vertex AI) · Databricks · MLflow · Docker · Git / GitHub CI/CD · FastAPI
Viz & BI	Tableau · Power BI · Streamlit · Plotly / Dash · Matplotlib / Seaborn
Statistics	A/B Testing & Experiment Design · Hypothesis Testing · Bayesian Inference · Causal Inference · Time Series Analysis

WORK EXPERIENCE

Current Company Name City, State (or Remote)
Data Analyst / Junior Data Scientist Aug 2024 – Present

- Built **multimodal image recognition system** using Gemini Vision API to dynamically classify and describe 500K+ product images, achieving **94.3% accuracy** and reducing manual tagging effort by 70%
- Engineered **GenAI orchestration pipeline** for enterprise library system using LangChain + Gemini API, automating document retrieval across 200K+ records and cutting search time from 8 min → **45 seconds**
- Developed automated reporting pipeline in Python + SQL processing **500K+ records monthly**, reducing report generation time by 60% and enabling real-time dashboards for 12-person leadership team
- Collaborated cross-functionally with product and engineering teams to translate business KPIs into measurable data models, directly influencing **3 product roadmap decisions**

Previous Company Name City, State
Data Analyst Jun 2021 – Jul 2024

- Analyzed **2M+ customer transaction records** using Python (Pandas, NumPy) and SQL to identify churn patterns, enabling targeted retention campaigns with **18% churn reduction**
- Built predictive model using **XGBoost** to forecast monthly revenue with **91% AUC-ROC**, deployed via Flask API and integrated into live business intelligence dashboards
- Designed and executed **A/B test** across 500K users to evaluate new email recommendation algorithm; validated **12% CTR lift** (p<0.001) driving \$420K incremental annual revenue
- Created executive-level **Tableau dashboards** consolidating 8 data sources; reduced weekly reporting overhead by 5 hrs/week across 4 teams
- Standardised ETL pipeline using **Apache Airflow**, reducing data freshness lag from 24h → **30 minutes** for 6 downstream business teams

PROJECTS

◆ **Real-Time Big Data Analytics Pipeline** github.com/yourname/bigdata-pipeline

Apache Kafka · PySpark Structured Streaming · Delta Lake · AWS S3 · Grafana

- Architected end-to-end streaming pipeline ingesting **2M+ events/day** via Kafka → PySpark → Delta Lake; achieved **sub-2s latency** for real-time analytics dashboard
- Reduced data storage costs by **40%** using Delta Lake's ACID transactions and intelligent partitioning on AWS S3

◆ **RAG-Powered Intelligent Document Search (Gemini + ChromaDB)** github.com/yourname/rag-library · [live demo](#)

Gemini API · LangChain · ChromaDB · FastAPI · Streamlit · Python

- Built production RAG pipeline over 200K+ documents; achieved **RAGAS relevance score 0.87**, reducing hallucination rate by 65% vs. naive LLM querying
- Deployed via FastAPI + Streamlit serving **1K+ daily queries**; wrote architectural deep-dive published on Medium (2.4K+ views)

◆ **Customer Churn Prediction — End-to-End ML System** github.com/yourname/churn-prediction

XGBoost · Scikit-learn Pipelines · MLflow · Docker · AWS SageMaker

- Engineered full ML pipeline (feature engineering → XGBoost → MLflow tracking → Docker → SageMaker); achieved **91% AUC-ROC** on 500K-row dataset

◆ **A/B Testing & Experimentation Framework** github.com/yourname/ab-testing-framework · [Streamlit app](#)

Python · SciPy · Bayesian A/B Testing · Streamlit · Plotly

- Built interactive Streamlit app implementing both frequentist and Bayesian A/B testing with **power analysis, MDE calculator, and sequential testing**; used internally by 3 teams

EDUCATION

Florida Atlantic University — Boca Raton, FL Graduated: [Month, Year]
Master of Science, Data Science and Analytics GPA: X.X / 4.0

Relevant Coursework: Machine Learning, Deep Learning, Statistical Inference, Big Data Analytics, Natural Language Processing, Data Mining, Time Series Analysis

CERTIFICATIONS & ACHIEVEMENTS

Big Data Analytics Certification	Florida Atlantic University	2024
Google Professional Data Engineer (In Progress / Completed)	Google Cloud	2025
Databricks Certified Associate Developer for Apache Spark	Databricks	2025
Kaggle Competition — [Competition Name] · Top X% (Bronze/Silver Medal)	Kaggle	2025

How to Customise This Resume for Maximum Impact

✔ Bullet Writing Formula

Formula: Action Verb + Tool/Method + What you did + Measurable result

Example: "Built XGBoost model to predict churn, achieving 91% AUC-ROC reducing churn 18%."

Always quantify: Use %, \$, records count, users, time saved, latency, accuracy

Never say: "Responsible for" or "Assisted with" — use strong action verbs

Action verbs: Built, Engineered, Designed, Developed, Deployed, Implemented, Automated, Reduced, Improved

🎯 ATS Optimisation Rules

Copy keywords verbatim from each job description into your resume

Use Jobscan.co to verify ATS score before sending — target 80%+

Standard headers only: Experience, Education, Skills, Projects, Certifications

No columns in the ATS version (save this visual version for human review)

Spell out acronyms once: "Natural Language Processing (NLP)"

🔥 Highlight Your GenAI Work

Lead with Gemini API + multimodal AI — it's your #1 differentiator in 2026

Upgrade library project to RAG pipeline before featuring it — adds massive depth

Include a GitHub link and live demo URL for every project

Add RAGAS scores, accuracy percentages, and latency numbers to every AI project

Mention "Multimodal AI" explicitly — it's a rare and hot keyword right now

Skills Section Strategy

Only list skills you can discuss confidently in an interview

Put Python and SQL first — they're in every DS job description

Group GenAI tools separately — it signals you're a modern DS candidate

Include both the library and the use case: "Scikit-learn Pipelines" not just "ML"

Remove skills you're weak at — they can become interview traps

Length & Format Rules

Strictly 1 page for 0–4 years experience — no exceptions for FAANG

Remove GPA if below 3.5, include if above

Education goes AFTER experience once you have work history

No photos, no personal pronouns, no "References available"

Use consistent date format throughout: "Jan 2024" or "2024" — not both

Company-Specific Tuning

Google/DeepMind: Emphasise stats, publications, research rigour

Meta: Lead with A/B testing, product metrics, experimentation at scale

Amazon: Frame bullets around Leadership Principles (impact, ownership, data-driven)

Netflix: Highlight recommender systems and A/B testing culture fit

Databricks/Snowflake: Push your Big Data cert and Spark/PySpark work hard

 **Download as PDF:** Click "**Download PDF**" in the toolbar above → Browser print dialog opens → Set **Destination = "Save as PDF"** → Paper size = **A4** → Disable headers & footers → Disable margins (or set to Minimum) → Save. This produces a clean, single-page, print-quality PDF that is identical to what you see above.